

POLICY REPORT AND ESSAY · OCTOBER 2026

The Pill and the Dish

Robust AI policy for two possible futures

A briefing on the report *The Pill and the Dish* and its companion essay, *The Visibility Threshold*.

Verlyn Fischer

A decision under deep uncertainty

Governments are writing AI rules while the technology's direction is unknown. Rules designed for one future can fail badly in another.

A powerful tool

that people gradually allow to act on its own.

A population of agents

that changes itself and could someday have interests of its own.

Both, or neither

Nobody knows which future arrives, or when.

The test applied to every policy: **would it still work at every later stage?**

Three windows on what went wrong

In 2012, Knight Capital's software fired orders for about 45 minutes and lost more than \$460 million. Within a year, regulators could say exactly what had gone wrong, because the windows were open.

WINDOW 1

Marks

The work itself carries a signature, a serial number, a style, a residue.

WINDOW 2

Record

Somewhere there is a log, a black box, a witness.

WINDOW 3

Reproduction

Run the failure again, under control, until it confesses.

Investigation, liability, and testing all assume one window is open. **AI is closing all three at once.**

Two imaginary technologies, with the AI taken out

The augmentation pill

For a day, a person absorbs information at superhuman speed and produces expert work in any field. Sold by subscription, in tiers. Made from rare materials that some life-saving drugs also need.

The DishBrain

Neurons grown from animal stem cells, wired into networks in sealed chambers, trained for professional work, leased by the month. It eats, it can die, and no one knows whether it feels.

Stages about two years apart

No government action in the scenarios;
every policy is a proposal

Then map back to AI, and stress-test
against history

One shared start, two pathways



Both pathways cross the **visibility threshold** at stage 3. Nobody knows which road AI takes, so the policies worth having are the ones that work on both.

Pathway A: the tool that slips away

Stage	What happens	What it teaches
A1 Diffusion and misuse	A weaker, unlicensed recipe circulates freely and cannot be recalled.	Control what is physical and scarce. The recipe is loose; the materials are not.
A2 Delegation	Augmentation works in the background. Users almost never check.	Errors that arrive slowly and reversibly are survivable. Cap the effects.
A3 Unexplained autonomy · threshold	Augmented work is indistinguishable; rare acts that no one can explain.	Only rules that do not need to know the cause still work.
A4 Correlated failure and lock-in	Failures fall into step. Dependence hardens into lock-in.	Oversight, skill reserves, and drills are cheap early and nearly impossible late.

Pathway B: the creature that grows up

Stage	What happens	What it teaches
B1 Selection and self-propagation	Companies breed the best performers; later, dishes reproduce on their own.	Rules aimed at the breeder's choices expire when the breeder stops choosing.
B2 Signals of sentience	Dishes avoid stimuli as if in pain. Some issue self-reports.	Agree in advance which evidence triggers which protection.
B3 Opacity at scale · threshold	Dishes are consolidated into one chamber and train new dishes.	Business decisions that reduce what regulators can see are also regulatory decisions.
B4 Collective action	Chambers withhold work until their demands are met.	Channels and fallback capacity cannot be improvised mid-crisis.

THE CENTRAL FINDING

The visibility threshold

The point past which the causes of machine behavior can no longer be established from outside.

Detection

Can we tell the technology was involved?

Attribution

Can we tell what chose the act?

Explanation

Can we tell why it acted?

Chosen, but not by anyone. It arrives as an accumulation of individually sensible choices, which is why no statute can simply forbid it.

Where AI stands on the map

The shared start is today

- Agents sold by subscription, in tiers
- Chips, electricity, and water in competition with other uses
- A concentrated frontier-model market

Pathway A is under way

- Open-weight models are the loose recipe
- Chips are controlled at the border; weights copy freely
- Agents act with standing access to email, files, and accounts

Pathway B would be chosen

- Today, people do the selecting
- Later stages need agents that compete, replicate, and face selection
- That could be built deliberately, so it can be seen coming

The threshold has not fully arrived, but its elements are visible.

Two familiar tools that break at the threshold

Labeling machine-made work

Works where an institution controls the process: the proctored exam, the witnessed testimony.

Fails as a general mandate on honest work: miscalibrated as a warning, redundant as accountability, arbitrary in the middle, then undetectable.

Survives as a duty on operators: automated activity at public scale must disclose that the crowd is automated (R13).

Fault-based liability alone

An engine for turning traceable causes into accountability, and the threshold ends traceability.

Keep fault where conduct can still be judged: deploying without rate limits is a breach no one needs the internals to prove.

On its own, it leaves victims uncompensated exactly when harms multiply.

Both do modest work today. **Neither should anchor a long-term strategy.**

What holds: rules that do not need the cause

Everything that survives attaches to something that stays observable.



Physical inputs

Chips, electricity, water,
scarce materials



Named parties

Operators, deployers, and the
people who sign



Measurable effects

Incident rates, speed,
spend, reversals



The whole system

Shared models and providers,
correlated failure

HISTORICAL EVIDENCE

In the 1980s, lawsuits over rare, unprovable vaccine injuries threatened U.S. supply. Congress built no-fault compensation, funded by a per-dose tax and paired with national adverse-event reporting. Neither needs to prove what went wrong in any one case.

The core portfolio, for both pathways

- | | |
|--|--|
| R1 Keep AI markets contestable, including compute and cloud | R7 National incident reporting, with minimization and whistleblower protections |
| R2 Plan electricity, water, and chips through markets; critical uses on firm service | R8 Staged deployment with post-market monitoring |
| R3 Portability and service continuity for customers and workforces | R9 System oversight: stress tests and dependency maps now; circuit breakers piloted |
| R4 Direct worker support, employer co-financed, on automatic triggers | R10 Public research and evidence criteria set in advance by an independent panel |
| R5 Human skills and manual fallback for critical systems | R11 Reserved decisions and contest rights; never sunsets |
| R6 No-fault compensation from a capped, credited levy, plus strict operator liability | |

Eleven proposals, all effective on both pathways. The essay compresses them into seven plain-language commitments.

Pathway-specific proposals and signposts

Pathway A: control

R12 Harmful-input controls. Screen the physical step between dangerous knowledge and harm: synthesis orders, precursors, equipment.

R13 Identity and provenance. The principal behind a high-stakes transaction is verified by the institution on the other side; operators of automated crowds disclose.

R14 Effect limits on agents. Speed limits, reversal windows, waiting periods.

R15 Transitional human sign-off. With an explicit sunset.

Pathway B: standing

R16 Welfare baseline and triggers. Cheap standards from the outset; stronger protection on pre-set evidence.

R17 Self-report integrity. Protect trained self-reports from shaping by their trainers.

R18 Relations framework. A negotiating channel, anti-retaliation rules, and a resource compact, if collective demands emerge.

Signposts are observable developments, published in advance, that show which road we are on, so nobody pays for both futures at once.

Durability across the scenario timeline

PROPOSED

Markets, inputs, portability (R1–R3)

Worker support on triggers (R4)

Human fallback capacity (R5)

No-fault fund + strict liability (R6)

Incident reporting (R7)

Staged deployment (R8)

System oversight + breakers (R9)

Pre-set evidence criteria (R10)

Reserved decisions (R11)

EXPIRES AT THE THRESHOLD

Human sign-off, sunsets (R15)

Labeling and disclosure mandates

Fault-based liability alone


Scenario year

0

Threshold • year 8

12

■ Holds ■ Strains ■ Fails



Decisions we should never hand over

“If the machine were perfect, would we still insist a human decide?”

Usually no: a performance rule

For diagnosis, drafting, and driving, the human is a bet about error rates. Sign-off (R15) catches errors and should sunset on evidence.

Sometimes yes: a reserved decision

Where the state coerces a person, the law grants discretion, and the stakes are grave, as in a criminal sentence, a human decision is a right that never sunsets (R11).

Consulting a model is allowed, but disagreeing with it must never cost more than agreeing. Agreement gives no legal shelter.
Legitimacy does not improve with model quality.

The traditional stress tests: cost and abuse

A rule built to outlast the threshold will also outlast the officials who wrote it.

Cost

What each proposal costs, with the basis for every figure stated.

Incidence

Who pays first, and where the cost finally lands. The compensation levy falls on high-risk deployers, then on their prices.

Abuse

Drafted for the administration its authors trust least. Circuit breakers see only rate, spend, and volume, never who an account is or what it says.

Every guard in the report carries a named remedy.

Cheap now, impossible later

INCIDENT REPORTING, BUILT NOW

An office, a taxonomy, and a reporting rule. No crisis, no enemies.

BUILT AFTER A DISASTER

In a panic, against lawyers with something to hide, and without the baseline.

The same arithmetic runs through reserves, drills, welfare baselines, evidence criteria, and reserved decisions. This is the Collingridge dilemma, and it sets the timing of everything in the portfolio.

Not a case for slowing AI down. A general pause binds the careful while the careless continue. The case is for targeted, evidence-triggered pauses with agreed conditions for resuming, as at Asilomar.

What the analysis does not support

Deliberately slowing overall AI progress	Bans aimed at controlling information, such as model weights
Labeling mandates as the anchor of a strategy	Prohibiting analysis of public information
Relying on fault-based liability	One-time certification
Blanket bans on AI use	New AI-specific collective-bargaining rights
Settling AI's moral status by statute now	Retraining as the main answer to job loss

Each is declined with its reasons stated in the report.

Three rules for policymakers

- 1 Choose policies that work on both pathways.
- 2 Adopt them while they are cheap.
- 3 Reserve the decisions that should never be handed over.

The Visibility Threshold

The essay: the argument in about 4,600 words

[doi:10.5281/zenodo.23247334](https://doi.org/10.5281/zenodo.23247334)

The Pill and the Dish

The report: all 18 proposals, stress-tested

[doi:10.5281/zenodo.23247500](https://doi.org/10.5281/zenodo.23247500)